

On determining the post-test probability of covid-19

Sam Weisenthal (firstnamelastname@gmail.com)

Contents

Summary	1
Background	2
Marginal	2
Bayes rule	2
Estimates from (Bruemmer et al. 2021)	3
Implementation	3
Confidence bounds (work in progress)	5
Difficulties	6
Plots using (Bruemmer et al. 2021)	6
Conditional (work in progress)	12
Estimates from (Prince-Guerra et al. 2021):	13
Other variables that might be relevant:	13
Weakening the assumption of negligibility (aka, that the marginal estimate is fine):	14
(Prince-Guerra et al. 2021) vs (Bruemmer et al. 2021)	14
On which bound the marginal should favor	15
Bound	19
Ideally	19
On the test cutoff (if there is a cutoff)	20
Serial testing	20
Viral load	20
Positive tests	20
Code link	21
References	21

Summary

Tests are crucial for controlling the covid-19 pandemic. Test result interpretation is also important. For example, a test could be negative, but covid-19 might still be present. After getting a negative rapid antigen test a few months ago, I tried to use a standard approach to interpret the result. Namely, I used Bayes rule to calculate post-test probability. This was not easy. I recorded my efforts at <https://discourse.datamethods.org/t/determining-post-test-probability-of-covid-19/4723/12>, and I have tried to collect my most updated thoughts at <https://github.com/samuelweisenthal/BinaxNOWpostTestProb/blob/main/covid.post.test.pdf>. Difficulties included specifying my pretest probability, getting appropriate estimates for the test operating

characteristics, and understanding my estimate in light of other information I had available, such as exposure and symptom status. Beware of the trap: your pretest probability cannot depend on variables that are different from the ones collected in the study that evaluated the test. (See also the diagnosis chapter in <http://hbiostat.org/doc/bbr.pdf>.) This trap is particularly problematic when you have reason to believe that certain variables, such as symptom status or exposure, make it more likely that you have the disease at baseline.

Overall, I found this to be a hard problem, based both on inherent difficulties and also on the indirect way in which I had to go about it. There may be ways to make the entire process more direct, shifting the burden of test interpretation from the person taking the test to statistical analysts who have appropriate tools and databases. One way might be to maintain databases that could be used to fit models for post-test probability that are specific to location and time, as well as other important covariates. In this case, someone would just need to go to a website and enter their information to get a test interpretation.

Background

Some months ago, I wanted to, given what I had available, try to determine my post-test probability of covid-19 after a negative Binax NOW rapid antigen test and a sustained, unmasked exposure. Originally, I hoped we could apply Bayes rule to some estimates from the literature. It turns out to be more complicated.

Marginal

Bayes rule

To begin this document, I will consider the easier problem of estimating post-test probability given a negative test only, with no exposure. In other words, let us forget about the conditional post-test probability, $p(dz+|test-, exposure)$, and consider the “marginal” $p(dz+|test-)$ (this is almost always done, unfortunately, when we talk about post-test probabilities). We have

$$\begin{aligned} p(dz+|test-) &= \frac{p(test-|dz+)}{p(test-)} p(dz+) \\ &= \frac{1 - p(test+|dz+)}{p(test-|dz-)p(dz-) + p(test-|dz+)p(dz+)} p(dz+) \\ &= \frac{1 - sens}{spec(1 - pretest) + (1 - sens)pretest} pretest \end{aligned}$$

On renaming

If we can fill in the pieces of this equation, we can provide an estimate of the left hand side. Note that in the last line, I rename probabilities as $sens=sensitivity$ and $spec=specificity$. This renaming, interestingly, can cause some difficulty; we understand the terms in the second to last line, but the renaming in the last line renames and strips the terms of their probabilistic meaning (along with removing information about what is being conditioned upon), which can lead to misuse, and may be related to the persistence of the framework for test interpretation that gives so many issues below (Moons and Harrell 2003). It is worth noting that test interpretation might be much better if there were databases kept such that we could just estimate $p(dz+|test-, other\ variables)$ directly, rather than having to use Bayes rule. This type of modeling does not occur usually because each model would be specific to time and location, at least. This sounds like a daunting task, therefore, to fit all of these models, but the reality is that if this is not done that way, it has to be done another way — it just shifts the responsibility of interpretation from the test manufacturer or some organization with appropriate expertise onto the person using the test, and we are left with nothing but the equation above, as far as I know.

Estimates from (Bruemmer et al. 2021)

Here, I am going to use this equation to calculate the post-test probability given a negative test for a battery of rapid antigen tests. Consider numbers from the recent meta-analysis (Bruemmer et al. 2021). Note that these estimates, as far as I know, condition on nothing except the test result [and adult status? read more carefully] (I do not know how they would otherwise pool so many estimates). The validity of the estimates depend on the designs of the studies that were analyzed in the meta-analyses. According to (Harrell 2020), a retrospective case-control is most appropriate. We assume that the studies used to obtain the sensitivities and specificities had a reasonable design, which is not always the case.

Implementation

I am going to show my code to promote transparency and reproducibility. You can find the notebook in my GitHub; use at your own risk. You are free to use and modify this code as long as you allow others to use and modify your modification of this code.

We start with the likelihood ratio.

```
like.rat=function(sens,spec,pretest){(1-sens)/((1-sens)*pretest+(spec)*(1-pretest))}
```

We then get that the post-test probability of disease given a negative test is likelihood ratio times the pretest probability.

```
post.test.prob = function(sens,spec,pretest){  
  like.rat(sens=sens,spec=spec,pretest=pretest)*pretest  
}
```

Let us compute the post-test probability over a grid of pre-test probabilities.

```
pre.tests = seq(0,1,0.01)  
get.post.tests = function(sens,spec,pre.tests){  
  post.tests = c()  
  for(i in 1:length(pre.tests)){  
    post.tests[i]=post.test.prob(sens=sens,spec=spec,pretest=pre.tests[i])  
  }  
  post.tests  
}
```

Now create a function to plot

```
plot.p = function(lns,post.label,xlab,  
                  leg.labels,main){  
  
  lwd.point = 2  
  plot(c(0,0),pch=0,xlim=c(0,1),ylim=c(0,1),  
        xlab=xlab,ylab=post.label,  
        lty=1,lwd=lwd.point,main=main,col='white')  
  
  xseq = seq(0,1,0.1)  
  yseq = seq(0,1,0.1)  
  for(x in xseq){  
    abline(v=x,col='gray')  
  }  
  
  for(y in yseq){  
    abline(h=y,col='gray')  
  }  
  colorBlindBlack8 <- c("#000000", "#E69F00", "#56B4E9", "#009E73",
```

```

                                "#F0E442", "#0072B2", "#D55E00", "#CC79A7")
nl = length(lns)
for (i in 1:nl){
  lines(pre.tests,lns[[i]]$m,lty=1,lwd=lwd.point,col=colorBlindBlack8[i])
  lines(pre.tests,lns[[i]]$u,lty=2,col=colorBlindBlack8[i])
  lines(pre.tests,lns[[i]]$l,lty=2,col=colorBlindBlack8[i])

  #rgbval = paste(as.vector(col2rgb(colorBlindBlack8[i])), collapse = " ")
  #lines(pre.tests,lns[[i]]$u,lty=2,col=rgb(rgbval[1],rgbval[2],rgbval[3]))
  #lines(pre.tests,lns[[i]]$l,lty=2,col=rgb(rgbval[1],rgbval[2],rgbval[3]))

}

legend("topleft",legend=leg.labels,bg='white',
      lty=rep(1,nl),col=colorBlindBlack8[c(1:nl)],lwd=c(lwd.point,lwd.point))
}

```

Here I have converted the table from (Bruemmer et al. 2021) into a list.

```

tests=list(
  veritor=list(
    nm='BD Veritor',
    sens=list(m=.64,l=.49,u=.76),
    spec=list(m=1,l=.99,u=1)
  ),
  Binax=list(
    nm='BinaxNOW',
    sens=list(m=.62,l=.48,u=.74),
    spec=list(m=1,l=1,u=1)
  ),
  Clinitest = list(
    nm='CLINITEST',
    sens=list(m=.62,l=.47,u=.75),
    spec=list(m=.99,l=.97,u=.99)
  ),
  Coris = list(
    nm="Coris",
    sens=list(m=.40,l=.29,u=.52),
    spec=list(m=.99,l=.95,u=1)
  ),
  LumipulseG = list(
    nm="Lumipulse G",
    sens=list(m=.87,l=.78,u=.93),
    spec=list(m=.97,l=.89,u=.99)
  ),
  LumiraDx = list(
    nm="LumiraDx",
    sens=list(m=.88,l=.59,u=.98),
    spec=list(m=.99,l=.96,u=1)
  ),
  Panbio = list(
    nm="Panbio",
    sens=list(m=.72,l=.65,u=.78),

```

```

    spec=list(m=.99,l=.99,u=1)
  ),
  Rapigen = list(
    nm="Rapigen",
    sens=list(m=.62,l=.47,u=.75),
    spec=list(m=.99,l=.94,u=1)
  ),
  Sofia = list(
    nm="Sofia",
    sens=list(m=.77,l=.74,u=.80),
    spec=list(m=.99,l=.98,u=1)
  ),
  StandardF = list(
    nm="StandardF",
    sens=list(m=.68,l=.56,u=.79),
    spec=list(m=.98,l=.97,u=.99)
  ),
  StandardQ = list(
    nm="StandardQ",
    sens=list(m=.75,l=.69,u=.80),
    spec=list(m=.99,l=.98,u=1)
  ),
  StandardQnasal = list(
    nm="StandardQnasal",
    sens=list(m=.80,l=.70,u=.87),
    spec=list(m=.99,l=.98,u=1)
  )
)

```

Confidence bounds (work in progress)

We want to take into account the uncertainty in the estimates above. We need to choose which values of sensitivity and specificity will be used for the confidence bounds. I am pretty sure I am correct about sensitivity here. I am not sure about specificity, but there is very little variation in specificity over the tests anyway. Recall from above that

$$p(dz + |test-) = \frac{1 - sens}{spec(1 - pretest) + (1 - sens)pretest} pretest$$

Note that this implies that sensitivity of 1 gives 0 post-test and sensitivity of 0 gives $\frac{pretest}{spec(1 - pretest)}$. So higher sensitivity at fixed specificity will decrease post-test probability of the disease being present.

Fixing sensitivity, $spec = 0$ implies $\frac{1 - sens}{sens}$ and $spec = 1$ implies $\frac{(1 - sens)}{(1 - pretest) + (sens)(pretest)} pretest$. This is not as clear, but I guess I will assume that higher specificity will decrease post-test probability since it is in the denominator in the original expression.

Hence, our lower bound on post-test probability will use the lower confidence bound of sensitivity and the lower confidence bound of specificity (and vice versa for the upper bound on post-test probability).

```

get.t.posts = function(tests){
  test.posts=list()
  for(t in tests){
    test.posts[[t$nm]]=list(
      m=get.post.tests(sens=t$sens$m,spec=t$spec$m,pre.tests),

```

```

    l=get.post.tests(sens=t$sens$l,spec=t$spec$l,pre.tests),
    u=get.post.tests(sens=t$sens$u,spec=t$spec$u,pre.tests)
  )
}
test.posts
}

```

Difficulties


We will plot now the post test probabilities for each test as a function of the pretest probability and the operating characteristics above. Note that these plots may be falsely comforting. First, how do we specify the pretest probability? It's not necessarily prevalence, because prevalence is conditional upon location and time, which are conditional upon the geographic and temporal characteristics of the population. However, the sensitivity and specificity estimates do not account for these things. Second, as I will discuss below, assuming I can specify a reasonable pretest probability, what if there are things about me that are important to take into consideration? I discuss exposure below, but how about symptom status, variant, time since symptom onset, etc? *You cannot simply change your pretest accordingly; sensitivity and specificity must also be updated.* To see this, look back at the derivation above, or read the diagnosis chapter in (Harrell 2021).



Plots using (Brummer et al. 2021)



```



#par(mfrow=c(12,1))
for (k in 1:12){
test.posts = get.t.posts(tests[k])
plot.p(test.posts,
  post.label="Post-test: p(dz+|test-)",xlab="Pre-test: p(dz+)",
  leg.labels=names(test.posts),main=""
)
}



```


















Conditional (work in progress)

Let us consider the general problem of conditioning on other covariates. Note, again, that the post-test estimates above do not take into account the extra variable that we call “exposure.” Before thinking about exposure, let us consider a different variable: symptoms. Define sx =symptoms, ad = adult, and dz =disease (covid-19). I am including symptom status and “adult” as done in a study that measures these things, (Prince-Guerra et al. 2021), but this is more for the sake of illustration than to give an accurate estimate.

Note:

$$\begin{aligned}
 p(dz + |test-, sx) &= \frac{p(test- | dz+, sx)}{p(test- | sx)} p(dz + | sx) \\
 &= \frac{1 - p(test+ | dz+, sx)}{p(test- | dz-, sx)p(dz- | sx) + p(test- | dz+, sx)p(dz+ | sx)} p(dz + | sx) \\
 &= \frac{1 - sens}{spec(1 - pretest) + (1 - sens)pretest} pretest
 \end{aligned}$$

Aside: will there be issues with the observational nature of the data? note that, just focusing on $p(dz+|test, sx)$, we are really interested in the test effect. Let us therefore consider how this would be estimated in general. First, if the data is observational and there is an unmeasured variable that affects whether someone takes a test and also whether they have the disease (confounding), the effect estimate may be incorrect. Generally, we would condition on the confounder to remove this issue. I am still trying to think more about how this relates to the discussion below on conditioning on the necessary set of variables. I think it is different, since we can defend a marginal (non-conditioned) post-test probability, if it is tailored enough to our covariates to inform decision-making, but we can’t really defend an unadjusted estimate of, e.g., a treatment effect. This actually seems to be the tension here. In some cases, marginal, $p(dz + |test-)$, is good enough, even though ideally we would have conditional $p(dz + |test-, age, sex, location, time, incubation, \dots)$. Especially, marginal is good enough if somehow the variables we do not condition on in the marginal estimate sort of cancel each other out.

Estimates from (Prince-Guerra et al. 2021):

Assuming that the study design of (Prince-Guerra et al. 2021) allows for valid estimation, we can use their point estimates and 95% confidence intervals for

$$p(test + |dz+, sx+, ad+) \approx 0.64 (0.57, 0.71)$$

and

$$p(test + |dz+, sx-, ad+) \approx 0.36 (0.27, 0.45).$$

Also,

$$p(dz - |test-, sx-, ad+) \approx 1(.99, 1)$$

and

$$p(dz - |test-, sx+, ad+) \approx 1(1, 1).$$

We can then plug these estimates into the equation above. Consider now what happens in the situation where instead of exposure, we did not measure symptom status:

$$p(dz + |test-) = (dz + |test-, sx+)p(sx + |test-) + p(dz + |test-, sx-)p(sx - |test-).$$

If we blindly compute the left hand side above, we will be estimating our post test probability ignoring symptoms, and therefore underestimate our risk. With the equation above, this is very clear that this should be the case. However, the mechanical use of Bayes rule to estimate post-test probability often leads us to ignore this fact, surprisingly. The same would be true for ignoring exposure; it may be the case that exposure is even more important than symptoms. This may also not be the case. Maybe, knowing that symptoms is more important than exposure, we can safely ignore exposure. Only gathering the appropriate data can tell us what is appropriate.

Other variables that might be relevant:

As I said above, we can only know how necessary it is to condition on additional covariates by collecting more data. However, one can also use previous studies to detect variables that might be relevant, but are ignored. I am concerned with the omission of gender, which appears to be possibly correlated with viral load (Mahallawi et al. 2021), assuming viral load and antigen presence are essentially the same. I think age should be treated as a continuous variable. It seems that age should correlate with viral load; this was not supported by Mahallawi, but it might be supported by (Euser et al. 2021). We need to condition on anything that leads to different antigen levels in different people. If the antigens are excreted or metabolized, we would need to take into account liver and kidney status. Of course, this should depend on immune system function, which will vary along with comorbidities and medications. Antigen level also probably depends on the covid strain. It appears that Table 2 in (Prince-Guerra et al. 2021) provides some interesting information that can help us here (there are counts of eg false positive and negative rates conditioned on sex, ethnicity, etc; unfortunately not jointly - I will have to look for the raw data). I am just eyeballing it, and it appears that the percentage of FN differs by sex, ethnicity, race, age, symptoms, days from symptom onset, and known exposure (!!).

I am also unsure how one should estimate $p(dz + |sx, ad)$. We can get this for the (Prince-Guerra et al. 2021) cohort, but this depends on things like location, time of year, and lockdown status. Some of these have changed a lot since the study was conducted. One can avoid having to think about the latter if things are done with updated, location-specific databases, as mentioned above. It really should not be the responsibility of the test taker to think about the pretest probability, especially because it depends on the covariates chosen in the study that evaluates the test. It is unreasonable to expect that someone can guess, without data, the value of their conditional pre-test probability.

Weakening the assumption of negligibility (aka, that the marginal estimate is fine):

I said before that possibly if the variables we do not condition on have equally opposing effect, we may be able to ignore them. This could allow us to relax the assumption that these variables are not necessary to adjust for. We really though have to gather data to see if this is the case.

Let us look at whether symptom status is reasonable to marginalize.

```
post.tests.6=get.post.tests(sens=0.64,spec=1,pre.tests)
post.tests.6l=get.post.tests(sens=0.57,spec=1,pre.tests)
post.tests.6u=get.post.tests(sens=0.71,spec=1,pre.tests)

l.6 = list(m=post.tests.6,l=post.tests.6l,u=post.tests.6u)

post.tests.36=get.post.tests(sens=0.36,spec=1,pre.tests)
post.tests.36u = get.post.tests(sens=.45,spec=.99,pre.tests)
#post.tests.36l = get.post.tests(sens=.27, spec=1,pre.tests)
post.tests.36l = get.post.tests(sens=.27, spec=1,pre.tests)

l.36 = list(m=post.tests.36,l=post.tests.36l,u=post.tests.36u)

# Prince-Guerra 2021

lns = list(l.6,l.36)
```


(Prince-Guerra et al. 2021) vs (Bruemmer et al. 2021)

We can compare $p(dz + |test-, sx, adult)$ from (Prince-Guerra et al. 2021) with $p(dz|test-)$ from (Bruemmer et al. 2021).

```
pts=get.t.posts(tests)
#pts$BinaxNOW
lns = list(l.6,l.36,pts$BinaxNOW)

plot.p(lns,
  post.label=" p(dz+|test-,...)",xlab="Pre-test: p(dz+|...)",
  leg.labels=c(TeX("sx+,adult+ (Prince-Guerra)"),TeX("sx-,adult+ (Prince-Geurra)"),"sx null, adult
  )
```

Prince-Guerra vs Brummer



On which bound the marginal should favor

It seems that the difference between the conditional and marginal estimates will be asymmetric. We expected that the marginal would be between the two conditionals; going back to α above, which combines the two conditional distributions, we see.

$$\begin{aligned} p(dz + |test-) &= p(dz + |test-, sx+)p(sx + |test-) + p(dz + |test-, sx-)p(sx - |test-) \\ &= p(dz + |test-, sx+)p(sx + |test-) + p(dz + |test-, sx-)(1 - p(sx + |test-)) \\ &= p(dz + |test-, sx+)\alpha + p(dz + |test-, sx-)(1 - \alpha), \end{aligned}$$

where $\alpha \in [0, 1]$.

Hence, the difference between the marginal estimate and your conditional estimate will vary depending on your covariates. In other words, for someone who is symptomatic, it might not be a big deal to use the marginal estimate. For someone who is asymptomatic, this might not be the case. I am fascinated that the marginal seems to favor the symptomatic curve; I would have expected the opposite. I am going to explore this a bit more.

We will make our variables $Z \sim \text{Bern}(\text{expit}(\beta X))$.

```
expit=function(x){exp(x)/(1+exp(x))}
#plot(-10:10, expit(-10:10))
```

Set the seed and sample size.

```
set.seed(1)
n=50000
```

Set prevalence, pdz .

```
pdz=.2
```

Let $dz \sim \text{Bern}(pdz)$, $sx \sim \text{Bern}(\text{expit}(-2 + (2+2)dz))$, so we generally don't have symptoms unless disease is present. Let $ts \sim \text{Bern}(\text{expit}(-3 + (3+0.5)dz + 1.5(sx)(dz)))$, so we tend not to have positive tests unless there is disease, and the presence of symptoms along with the disease increases the probability.

```
dz = rbinom(n,1,pdz)
sx=rbinom(n,1,expit(-2+(2+2)*dz))
#antigen = rnorm(n,mean=-3+sx*.1+dz*1) #todo: explore cutoff
ts = rbinom(n,1,expit(-3+(3+0.5)*dz+1.5*sx*dz))
df = data.frame(dz=dz,sx=sx,ts=ts)
```

```
print(c("p(dz+)",pdz))
```

```
## [1] "p(dz+)" "0.2"
```

```
pdzpsxn = dim(df[(df$dz==1)&(df$sx==0),])[1]/dim(df[(df$sx==0),])[1]
pdzpsxp = dim(df[(df$dz==1)&(df$sx==1),])[1]/dim(df[(df$sx==1),])[1]
print(c("p(dz+|sx-)",pdzpsxn))
```

```
## [1] "p(dz+|sx-)" "0.0330934856640668"
```

```
print(c("p(dz+|sx+)",pdzpsxp))
```

```
## [1] "p(dz+|sx+)" "0.646305563732705"
```

Calculate the marginal sensitivity and specificity

```
ptspdzp=dim(df[(df$dz==1)&(df$ts==1),])[1]/dim(df[(df$dz==1),])[1]
ptsndzn=dim(df[(df$dz==0)&(df$ts==0),])[1]/dim(df[(df$dz==0),])[1]
print(c("marg sens: p(test+|dz+)",round(ptspdzp,2)))
```

```
## [1] "marg sens: p(test+|dz+)" "0.85"
```

```
print(c("marg spec: p(test-|dz-)",round(ptsndzn,2)))
```

```
## [1] "marg spec: p(test-|dz-)" "0.95"
```

```
pdzptsn = dim(df[(df$dz==1)&(df$ts==0),])[1]/dim(df[(df$ts==0),])[1]
print(c("p(dz+|test-)",pdzptsn))
```

```
## [1] "p(dz+|test-)" "0.0374678306504516"
```

Calculate the conditional sensitivities and specificities. First, given symptoms

```
ptspdpsxp = dim(df[(df$ts==1)&(df$dz==1)&(df$sx==1),])[1]/dim(df[(df$dz==1)&(df$sx==1),])[1]
ptsndznsxp = dim(df[(df$ts==0)&(df$dz==0)&(df$sx==1),])[1]/dim(df[(df$dz==0)&(df$sx==1),])[1]
print(c("Cond sens, sx-: p(test+|dz+,sx+)",round(ptspdpsxp,2)))
```

```
## [1] "Cond sens, sx-: p(test+|dz+,sx+)" "0.88"
```

```
print(c("Cond sens, sx-:p(test-|dz-,sx+)",round(ptsndznsxp,2)))
```

```
## [1] "Cond sens, sx-:p(test-|dz-,sx+)" "0.96"
```

```
pdzptsnsxp = dim(df[(df$ts==0)&(df$dz==1)&(df$sx==1),])[1]/dim(df[(df$ts==0)&(df$sx==1),])[1]
print(c("p(dz+|test-,sx+)",pdzptsnsxp))
```

```
## [1] "p(dz+|test-,sx+)" "0.180345380096137"
```

Now no symptoms

```
ptspdpsxn = dim(df[(df$ts==1)&(df$dz==1)&(df$sx==0),])[1]/dim(df[(df$dz==1)&(df$sx==0),])[1]
ptsndznsxn = dim(df[(df$ts==0)&(df$dz==0)&(df$sx==0),])[1]/dim(df[(df$dz==0)&(df$sx==0),])[1]
print(c("Cond spec, sx+: p(test+|dz+,sx-)",round(ptspdpsxn,2)))
```



```
## [1] "Cond spec, sx+: p(test+|dz+,sx-) " 0.61"
print(c("Cond spec, sx+: p(test+|dz-,sx-)",round(ptsndznsxn,2)))

## [1] "Cond spec, sx+: p(test+|dz-,sx-) " 0.95"
pdzptsnsxn = dim(df[(df$ts==0)&(df$dz==1)&(df$sx==0),])[1]/dim(df[(df$ts==0)&(df$sx==0),])[1]
print(c("p(dz+|test-,sx-)",pdzptsnsxn))

## [1] "p(dz+|test-,sx-) " 0.0138754152335597"
Now look at  $P(sx|ts)$ 
psxptsn = dim(df[(df$sx==1)&(df$ts==0),])[1]/dim(df[(df$ts==0),])[1]
psxntsn = dim(df[(df$sx==0)&(df$ts==0),])[1]/dim(df[(df$ts==0),])[1]
print(c("P(sx+|test-)",round(psxptsn,2)))

## [1] "P(sx+|test-) " 0.14"
print(c("P(sx-|test-)",round(psxntsn,2)))

## [1] "P(sx-|test-) " 0.86"
It seems that it must favor the asymptomatic conditional. What is going on here? Let's check everything.
print(c("p(dz+|test-)",pdzptsn))

## [1] "p(dz+|test-) " 0.0374678306504516"
print(c("p(dz+|test-,sx+)p(sx+|test-)+p(dz+|test-,sx-)p(sx-|test-)",pdzptsnsxp*psxptsn + pdzptsnsxn*psxntsn))

## [1] "p(dz+|test-,sx+)p(sx+|test-)+p(dz+|test-,sx-)p(sx-|test-)"
## [2] "0.0374678306504516"
print(c("p(dz+|test-,sx+)",pdzptsnsxp))

## [1] "p(dz+|test-,sx+) " 0.180345380096137"
print(c("p(dz+|test-,sx+)",get.post.tests(sens=ptspdzsxp,spec=ptsndznsxp,pre.tests=c(pdzpsxp))))

## [1] "p(dz+|test-,sx+) " 0.180345380096137"
#get.post.tests()
print(c("p(dz+|test-,sx-)",pdzptsnsxn))

## [1] "p(dz+|test-,sx-) " 0.0138754152335597"
print(c("p(dz+|test-,sx-)",get.post.tests(sens=ptspdzsxn,spec=ptsndznsxn,pre.tests=c(pdzpsxn))))

## [1] "p(dz+|test-,sx-) " 0.0138754152335597"
#cbind(pts$`Conditional sx+`$m,pts$`Conditional sx-`$m,pts$Marginal$m,pre.tests)[21,]
#cbind(pdzptsnsxp,pdzptsnsxn)

tests=list(
condsxp=list(
nm='Conditional sx+',
sens=list(m=ptspdzsxp,l=ptspdzsxp,u=ptspdzsxp),
spec=list(m=ptsndznsxp,l=ptsndznsxp,u=ptsndznsxp)
),
marg=list(
nm='Marginal',
```

```


sens=list(m=ptspdzp,l=ptspdzp,u=ptspdzp),
spec=list(m=ptsndzn,l=ptsndzn,u=ptsndzn)
),
condsxn=list(
  nm='Conditional sx-',
  sens=list(m=ptspdzp, l=ptspdzp, u=ptspdzp),
  spec=list(m=ptsndzn, l=ptsndzn, u=ptsndzn)
)
)

pts=get.t.posts(tests)

plot.p(pts,
  post.label="p(dz+|test-,...)",xlab="Pre-test: p(dz+|...)",
  leg.labels=c(TeX("Conditional: p(dz+|test-,sx+)"),"Marginal: p(dz+|test-)",TeX("Conditional: p(dz+|test-,sx-)"))
)

```

Conditional vs. marginal (simulation)



```

#cbind(pts$`Conditional sx+`$m*psxptsn + pts$`Conditional sx-`$m*psxptsn,
#pts$Marginal$m)

```

Again, this is contrary to my expectation. I will have to think more about this. I am pretty sure the following probability statement, copied from above, is not wrong...

$$\begin{aligned}
 p(dz + |test-) &= p(dz + |test-, sx+)p(sx + |test-) + p(dz + |test-, sx-)p(sx - |test-) \\
 &= p(dz + |test-, sx+)p(sx + |test-) + p(dz + |test-, sx-)(1 - p(sx + |test-)) \\
 &= p(dz + |test-, sx+)\alpha + p(dz + |test-, sx-)(1 - \alpha),
 \end{aligned}$$

In any case, it should be certain that the marginal falls within the conditionals. It should be certain that the marginal might be closer to one than to another. It is therefore the case that how much you will lose out by

constraining yourself to the study covariates depends on your covariates.

Bound

Recall that my original question was whether I had the disease given a sustained exposure. Unfortunately, (Prince-Guerra et al. 2021) does not give us operating characteristics that depend on exposure. However, the following section is just some musings on possibly being able to make a statement like “disease probability is greater than this value, even though we know it is not this value precisely.” We can say, for example, $p(dz + |test-) \leq p(dz + |test-, exposure+)$, which is based on the fact that a sustained exposure increases probability of disease.

We can also somewhat more circuitously write

$$p(dz + |test-, sx, ad, exposure+) = \frac{p(exposure + |test-, dz+, sx, ad)}{p(exposure + |test-, sx, ad)} p(dz + |sx, ad, test-)$$

We see that if

$$\frac{p(exposure + |test-, dz+, sx, ad)}{p(exposure + |test-, sx, ad)} > 1,$$

which is the case when the event “exposure” increases the probability of the disease (I assume this is true), then

$$p(dz + |sx, ad, test-),$$

which we can obtain from (Prince-Guerra et al. 2021), is a lower bound on

$$p(dz + |exposure+, sx, ad, test-).$$

So, even though the estimates in (Prince-Guerra et al. 2021) do not apply to our situation directly, if we were to assume that other variables, besides those concerning age cutoff, symptom status, test result, and exposure, are negligible, our post-test probability estimate can still be considered a lower bound, which still gives us some information, and also shows us, even with our strong assumptions, how little we actually know.

Ideally

Note that to obtain the graphs above, we are working with what we have available. This all gets somewhat confusing. One way to remove some (not all) difficulty here is to just take a direct approach from the beginning, which I will call the “ideal” approach. However, the following would be ideal: every time one takes a test, one goes to a website and enters information such as age, sex, zip code, and test result. The website then calculates post-test probability based on a model for, explicitly,

$$p(dz + |test, age, sex, symptoms, location, time, incubation, strain, ..)$$

The database and corresponding model would be updated in real-time based on geographically and temporally relevant statistics. Note that this post-test probability would depend very much on time, and therefore the model would have to be updated probably each day. It seems though that, barring sampling issues, we have this data. There is observational data collected when people report their test results (in other words, we collect information such as age, zip code, etc). Often, also, we have a PCR confirmation. It is unclear however whether we will have enough of the variables mentioned above, which are still necessary.

Note that age is given to us as “adult” - a large loss of info, symptom status is binarized, test result is binarized, and exposure was binarized by me, unfortunately (although we could just treat it as continuous). Also, ideally, the test result would be a continuous variable (eg, amount of viral load). This may be difficult due to at-home testing kit constraints. However, it seems currently that there is some cutoff above which a test is called positive. If it is possible for the tests to convey more information, such as through color or some type of numerical scale, it would lead to better estimates of the post-test probability (assuming there is no real hard cutoff - I am not sure).

On the test cutoff (if there is a cutoff)

It is not clear this is how it works, but, in general, test cutoffs have highly significant implications. If indeed there is a cutoff, what is the reward function that is being optimized? It appears that these tests were designed to minimize false positive results. However, that is not, in general, always a good idea. Decreasing false positives (e.g., by setting a high cutoff) also increases false negatives. In general, for someone who works in a highly populated area, or with vulnerable populations, a false negative is worse than a false positive. Therefore, it should be a priority not to include a threshold in the test, if possible.

As a thought experiment, consider a world in which people do not interpret tests, and just act according to the test result. In other words, if I test negative, I assume I have no disease, and if I test positive, I assume I have the disease. In this case, I believe (not sure) that the cutoff will set the reward that this society achieves entirely. In other words, it is important to interpret a test in order to avoid having your reward function be set externally by, e.g., a test cutoff. As we have seen, interpretation is complicated. Therefore, it is not unlikely that many people do act as if the test result is the true result, and therefore if there is a cutoff in the test, it has broad implications.

Serial testing

I have also done some more thinking on serial testing - my current thinking (maybe this is not correct, I need to still write it out here) is that **if the tests are independent**, you can essentially treat the post-test probability from the first test as the pre-test probability for the second. If this is the case, then two tests taken, premeditated, in sequence, will perform like a better test. If this is the case, using the plot, you can just start with a pretest and update it, and then make your next pretest your post-test from the first test.

Generally, though, it is also advised to take the two tests e.g. 24 hours apart to see if the viral load increases during that time. I am not sure that two tests that are taken like this are still independent.

Viral load

I mentioned viral load above. This was mentioned in a blog post <https://www.lesswrong.com/posts/Sct2SNhByS8Lzenub/rapid-antigen-tests-for-covid>. Viral load (obviously) affects sensitivity and specificity. It would be important to measure the time of supposed exposure in order to take this into account $p(dz + |..., \text{exposure time}, ...)$. Kevin Murphy also posted a good chart <https://twitter.com/sirbayes/status/1482522700195135496>.

Positive tests

Note that I am focusing only on the negative test case, although you could do the same for positive tests (I said originally that this was a non-issue, but I should not have — you can, of course, have a positive test with no disease, and I should repeat the analysis above for that case).

Code link

Code: <https://github.com/samuelweisenthal/BinaxNOWpostTestProb>

I appreciated comments and encouragement from the datamethods <https://discourse.datamethods.org/t/determining-post-test-probability-of-covid-19/4723/12> forum participants, Anna Park, and my brother on this topic. Please, if you have any comments, let me know!

References

Bruemmer, Lukas E, Stephan Katzenschlager, Mary Gaeddert, Christian Erdmann, Stephani Schmitz, Marc Bota, Maurizio Grilli, et al. 2021. “The Accuracy of Novel Antigen Rapid Diagnostics for Sars-Cov-2: A Living Systematic Review and Meta-Analysis.” *medRxiv*. Cold Spring Harbor Laboratory Press.

Euser, Sjoerd, Sem Aronson, Irene Manders, Steven van Lelyveld, Bjorn Herpers, Jan Sinnige, Jayant Kalpoe, et al. 2021. “SARS-Cov-2 Viral Load Distribution Reveals That Viral Loads Increase with Age: A Retrospective Cross-Sectional Cohort Study.” *medRxiv*. Cold Spring Harbor Laboratory Press.

Harrell, Frank E. 2020. <https://www.fharrell.com/post/backwards-probs/>.

———. 2021. <http://hbiostat.org/doc/bbr.pdf>.

Mahallawi, Waleed H, Ali Dakhilallah Alsamiri, Alaa Faisal Dabbour, Hamdah Alsaeedi, and Abdulmohsen H Al-Zalabani. 2021. “Association of Viral Load in Sars-Cov-2 Patients with Age and Gender.” *Frontiers in Medicine* 8. Frontiers: 39.

Moons, Karel GM, and Frank E Harrell. 2003. “Sensitivity and Specificity Should Be de-Emphasized in Diagnostic Accuracy Studies.” *Academic Radiology* 10 (6). Elsevier: 670–72.

Prince-Guerra, Jessica L, Olivia Almendares, Leisha D Nolen, Jayleen KL Gunn, Ariella P Dale, Sean A Buono, Molly Deutsch-Feldman, et al. 2021. “Evaluation of Abbott Binaxnow Rapid Antigen Test for Sars-Cov-2 Infection at Two Community-Based Testing Sites—Pima County, Arizona, November 3–17, 2020.” *Morbidity and Mortality Weekly Report* 70 (3). Centers for Disease Control; Prevention: 100.